

IT5005 Project Final Report (Milestone 4)

YOLO and Tracking Algorithms

Group 16

Team Members:

JIANG Zhuosong	A0314591E
LEE Bing Cheng	A0297134Y
LI Yizhe	A0304167J
LIU Yichao	A0304386A
MA Kuangxuan	A0304242W

April 18, 2025

Contents

1	Introduction	2
2	The Evolution of YOLO Architectures	2
2.1	Early Development: YOLOv1-v3	2
2.2	Middle Generation: YOLOv4-v8	2
2.3	State-of-the-Art: YOLOv9-v12	2
3	Architectural Components and Innovations	2
3.1	Attention Mechanisms in YOLOv12	2
3.2	Residual ELAN Blocks	2
3.3	Dynamic Architecture Configuration	2
3.4	C2f Backbone Module	3
3.5	Anchor-Free Detection Head	3
3.6	Decoupled Head Architecture	3
3.7	Loss Function Formulation	3
3.8	Dynamic NMS Implementation	4
4	Multi-Object Tracking Algorithms	4
4.1	Evolution of Tracking Algorithms	4
4.2	ByteTrack: Architecture and Innovations	4
4.2.1	Dual-Threshold Association Strategy	4
4.2.2	Motion Prediction with Kalman Filtering	4
4.3	ByteTrack Workflow and Advantages	5
5	Challenges and Advancements in Object Tracking	5
5.1	Integration of YOLO and Tracking Systems	5
5.2	Technical Challenges in Multi-Object Tracking	5
5.3	Performance-Resource Trade-offs	5
5.4	Data-Centric Improvements	5
5.5	Future Directions	5
6	Appendix: Figs & Tables	6
7	References	8

1 Introduction

Object detection evolved from traditional methods to deep learning approaches, with **YOLO (You Only Look Once)** addressing speed limitations by reformulating detection as a single regression problem. YOLO achieves real-time performance while maintaining competitive accuracy by predicting bounding boxes and classes in one evaluation. When combined with tracking algorithms like **ByteTrack**, YOLO enables real-time multi-object tracking across video frames, serving applications from autonomous vehicles to surveillance systems. This combination provides temporal consistency while reducing computation, making it ideal for time-sensitive applications (Terven et al., 2023).

2 The Evolution of YOLO Architectures

2.1 Early Development: YOLOv1-v3

YOLOv1 pioneered end-to-end detection with grid-based prediction, though limited in small object recognition. Subsequent versions introduced anchor boxes (YOLOv2) and residual connections (YOLOv3) for improved multi-scale detection.

2.2 Middle Generation: YOLOv4-v8

This phase established key architectural patterns including CSPDarknet backbones (YOLOv4) and PyTorch-based implementations (YOLOv5). Later versions introduced efficient rep structures (YOLOv6), enhanced feat-transfer modules (YOLOv7), and the lighter C2f blocks with anchor-free detection in YOLOv8 (Fig 1).

2.3 State-of-the-Art: YOLOv9-v12

Recent advancements focus on computational efficiency and attention mechanisms. YOLOv9 introduced Programmable Gradient Information (PGI) for reduced redundancy, while YOLOv10 optimized feature extraction through Redistribution of Intermediary Redundancy (RIR). YOLOv12 (Tian et al., 2025) marks a paradigm shift with its attention-centric design, addressing traditional vision transformer limitations with three key innovations: **Area Attention** reduces computational complexity to $\frac{1}{2}n^2hd$ by limiting attention to $\frac{1}{4}$ of the feature map **Residual ELAN (R-ELAN)** maintains feature integration while enabling effective backpropagation **Flash Attention** optimizes memory access through I/O reduction As shown in Fig 3, YOLOv12 achieves 58.9 AP on COCO with 0.29ms latency, outperforming previous versions through its hybrid transformer-CNN architecture and advanced distillation techniques.

3 Architectural Components and Innovations

3.1 Attention Mechanisms in YOLOv12

YOLOv12’s area attention mechanism operates as:

$$Attention(Q, K, V) = Softmax(\frac{Q_{sub}K_{sub}^T}{\sqrt{d}})V_{sub} \quad (1)$$

Where Q, K, V are subsampled to $\frac{1}{4}$ resolution. Combined with Flash Attention’s memory optimization (Dao et al., 2022), this achieves $3\times$ faster processing than standard attention implementations.

3.2 Residual ELAN Blocks

The improved R-ELAN structure enhances gradient flow through residual connections:

$$R-ELAN(X) = X + \sum_{i=1}^n Conv_{3\times 3}(Attn(Conv_{1\times 1}(X))) \quad (2)$$

This maintains 92

3.3 Dynamic Architecture Configuration

YOLOv12 adapts its architecture based on input complexity through:

$$Block_{active} = \begin{cases} R-ELAN & \text{if } \mathcal{H}(F) > \tau \\ CSP & \text{otherwise} \end{cases} \quad (3)$$

Where measures feature map entropy, dynamically choosing between attention and convolutional blocks. The technical innovations culminate in YOLOv12’s state-of-the-art performance (Table 1), achieving 63.1 AP on dense object detection tasks while maintaining real-time speeds under 2ms per image. This represents a 17% improvement over YOLOv8 with comparable computational resources, demonstrating the effectiveness of its attention-optimized architecture.

3.4 C2f Backbone Module

YOLOv8 replaces traditional CSP blocks with the C2f module, which optimizes computational efficiency. The C2f module can be mathematically represented as:

$$C2f(X) = Concat[X_{first}, F_n(F_{n-1}(\dots F_1(X_{second})))]) \quad (4)$$

Where X is split into X_{first} and X_{second} along the channel dimension, and F_i represents the i -th convolutional block in the module. Unlike standard CSP, C2f employs dense connections where each layer receives feature maps from all preceding layers:

$$F_i(X) = Conv(Concat[X, F_{i-1}, F_{i-2}, \dots, F_1]) \quad (5)$$

This formulation allows for more efficient gradient flow during backpropagation and reduces computational complexity by approximately 15% compared to YOLOv7’s backbone.

3.5 Anchor-Free Detection Head

YOLOv8 employs a fully anchor-free detection approach. For an input image divided into an $S \times S$ grid, each grid cell directly predicts:

$$\hat{y} = \{p_c, b_x, b_y, b_w, b_h, p_1, p_2, \dots, p_C\} \quad (6)$$

Where:

- p_c is the confidence score
- (b_x, b_y) are the center coordinates relative to the grid cell
- (b_w, b_h) are width and height relative to the input image
- p_i is the probability of class i given an object is present

Unlike anchor-based approaches, YOLOv8’s bounding box parameters are predicted as direct offsets from grid cell positions rather than as offsets from predefined anchor boxes.

3.6 Decoupled Head Architecture

YOLOv8 implements separate branches for classification, regression, and confidence estimation:

$$Classification : C(F) = \sigma(W_c \cdot F + b_c) \quad (7)$$

$$Regression : R(F) = W_r \cdot F + b_r \quad (8)$$

$$Confidence : O(F) = \sigma(W_o \cdot F + b_o) \quad (9)$$

Where F represents features from the neck, W and b are learnable parameters, and σ is the sigmoid activation function. This decoupling allows each task to optimize independently, reducing training conflicts.

3.7 Loss Function Formulation

YOLOv8 employs a composite loss function:

$$L_{total} = \lambda_{cls} L_{cls} + \lambda_{box} L_{box} + \lambda_{dfl} L_{dfl} \quad (10)$$

Where:

- L_{cls} is the Binary Cross-Entropy (BCE) loss for classification
- L_{box} is the CIoU loss for bounding box regression
- L_{dfl} is the Distribution Focal Loss

The CIoU loss is defined as:

$$L_{CIoU} = 1 - IoU + \frac{\rho^2(b, b^{gt})}{c^2} + \alpha v \quad (11)$$

Where ρ is the Euclidean distance between box centers, c is the diagonal length of the smallest enclosing box, v measures aspect ratio consistency, and α is a trade-off parameter.

The Distribution Focal Loss (DFL) models bounding box coordinates as continuous distributions rather than point estimates:

$$L_{dfl} = - \sum_{i=0}^n y_i \log(p_i) + (1 - y_i) \log(1 - p_i) \quad (12)$$

Where p_i represents the probability that the target falls into the i -th bin of a discretized range, and y_i is the ground truth distribution.

3.8 Dynamic NMS Implementation

YOLOv8 incorporates Dynamic NMS which adaptively adjusts the IoU threshold T based on detection density:

$$T = T_{base} \cdot \exp(-\beta \cdot d) \quad (13)$$

Where d is the local detection density, T_{base} is the base threshold, and β is a scaling factor. This approach more effectively handles crowded scenes by allowing more detections in dense regions while maintaining strict filtering in sparse areas.

These technical innovations collectively enable YOLOv8 to achieve 53.9 AP on COCO with an inference time of 0.39 ms on modern GPUs, representing a significant advancement in the speed-accuracy trade-off for real-time object detection.

4 Multi-Object Tracking Algorithms

4.1 Evolution of Tracking Algorithms

Multi-object tracking (MOT) algorithms have evolved from simple motion-based approaches to sophisticated systems integrating detection confidence, appearance features, and motion prediction. Early algorithms like SORT (Simple Online and Realtime Tracking) relied exclusively on Kalman filtering and IoU-based matching, offering speed but suffering from frequent ID switches. DeepSORT enhanced this by incorporating appearance features through ReID networks, improving tracking persistence at the cost of computational efficiency.

4.2 ByteTrack: Architecture and Innovations

ByteTrack (Zhang et al., 2022) represents a paradigm shift in tracking-by-detection approaches. Unlike previous methods that discard low-confidence detections, ByteTrack’s core innovation lies in its hierarchical association strategy.

4.2.1 Dual-Threshold Association Strategy

As shown in Fig 2, ByteTrack employs a two-stage matching process that intelligently utilizes all detection information:

1. **First Association:** Matches high-confidence detections ($score > thresh_{high}$) with existing tracks using IoU similarity and Kalman filter predictions.
2. **Second Association:** Recovers potentially valid objects from low-confidence detections ($thresh_{low} < score < thresh_{high}$) by matching them with previously unmatched tracks.

This approach addresses a fundamental limitation in tracking systems: the trade-off between precision and recall. By intelligently utilizing low-confidence detections that would otherwise be discarded, ByteTrack significantly improves tracking performance in occlusion scenarios.

4.2.2 Motion Prediction with Kalman Filtering

ByteTrack employs Kalman filtering for motion prediction, mathematically expressed as:

$$\hat{x}_t = Fx_{t-1} + Bu_t \quad (14)$$

$$P_t = FP_{t-1}F^T + Q \quad (15)$$

Where x_t represents the state vector (position and velocity), F is the state transition matrix, P_t is the covariance matrix, and Q represents process noise. This recursive estimation allows ByteTrack to predict object positions even during brief occlusions.

4.3 ByteTrack Workflow and Advantages

1. **Dual-Threshold Detection Association** (76.3 MOTA): Implements confidence-based detection grouping with high/low thresholds for staged trajectory matching
2. **Dynamic State Estimation**: Utilizes Kalman filtering for cross-frame prediction, reducing trajectory fragmentation by 32
3. **Robust Track Management**: Employs two-stage matching (high→low confidence) to achieve:
 - Low-confidence detection reuse: Improves recall by 12.6
 - Reduced ID switches: Maintains <2.1

Key technical features:

- Computational efficiency: 62 FPS on Tesla V100
- Architecture-agnostic design: Compatible with YOLO variants (v5-v12)
- Occlusion resilience: 89

This optimized implementation achieves state-of-the-art tracking performance while maintaining real-time processing capabilities.

As shown in Table 1, ByteTrack achieves the highest performance metrics while maintaining real-time capability, making it the preferred choice for applications ranging from autonomous driving to intelligent surveillance systems. A more detailed comparison of tracking algorithms is provided in Table 2.

5 Challenges and Advancements in Object Tracking

5.1 Integration of YOLO and Tracking Systems

Building on our discussion of YOLO architectures and ByteTrack, their integration forms a complete Tracking-by-Detection pipeline where YOLO provides detection data and ByteTrack associates objects across frames. This synergy leverages both the computational efficiency of YOLOv8’s C2f module and ByteTrack’s dual-threshold association strategy.

5.2 Technical Challenges in Multi-Object Tracking

Several challenges affect real-time tracking performance:

- **Object Scale Variation**: As objects move away from the camera, their size changes, affecting detection reliability—directly impacting YOLOv8’s anchor-free detection capabilities.
- **Small Object Detection**: YOLO architectures struggle with small objects due to limited feature representation, creating a weaker training signal (Kondrackis, 2024). This challenges the multi-scale detection capabilities discussed in Section 2.
- **Occlusion Handling**: Partially visible objects cause detection inconsistencies (Han et al., 2023), though ByteTrack’s second association stage helps mitigate this issue.
- **ID Switching**: Despite ByteTrack’s improvements, tracking algorithms still struggle with maintaining consistent object identity during brief occlusions (Huang et al., 2024).

5.3 Performance-Resource Trade-offs

As demonstrated in the evolution of YOLO architectures (Section 2), improving detection often means scaling computational resources:

- Larger YOLO models (YOLOv8x vs. YOLOv8n) offer better accuracy but at the cost of inference speed.
- This trade-off directly affects the real-time capability of the tracking system, particularly relevant to the architectural optimizations outlined in Section 3.

5.4 Data-Centric Improvements

Beyond algorithmic advances, dataset quality significantly impacts tracking performance (Famili et al., 1997). Well-curated training data that matches deployment conditions can substantially improve detection reliability before tracking algorithms are applied.

5.5 Future Directions

While this report focuses on Tracking-by-Detection with YOLO and ByteTrack, emerging end-to-end approaches like MOTRv2 (Zhang et al., 2023) are challenging the traditional paradigm. However, as evidenced by our architectural analysis, the optimal approach remains application-dependent, with speed-critical scenarios favoring the YOLO+ByteTrack combination discussed throughout this report.

6 Appendix: Figs & Tables

Algorithm	MOTA	IDF1	Real-time?	Key Feature
SORT	59.8	53.8	Yes	IoU + Kalman Filter
DeepSORT	61.4	62.2	No	Appearance + Motion
FairMOT	73.7	72.3	Yes	Joint Detection + ReID
OC-SORT	75.5	75.8	Yes	Camera Motion Compensation
ByteTrack	76.3	77.3	Yes	Dual-threshold Association

Table 1: Performance comparison on MOT17 benchmark

Algorithm	Key Feature	Strength	Weakness
SORT (2016)	Kalman + Hungarian Algorithm (IoU based)	Fast, simple to implement	Prone to ID switches, no appearance modeling
DeepSORT (2017)	ReID + Motion (Kalman Filter)	Robust to occlusion, more stable ID tracking, reduces ID switches by combining motion and appearance metrics	Computationally expensive due to feature extraction
FairMOT (2020)	Joint detection + ReID	Balanced detection and tracking accuracy in crowded scenes	Complex training
ByteTrack (2021)	Hierarchical association using high/low confidence detections	Recovers occluded or low-score objects	Relies on detector quality
OC-SORT (2022)	Motion-centric updates	Handles camera motion	Requires tuning for scenes, lacks appearance modeling

Table 2: Comparison of Tracking Algorithms

Layer Type	Complexity per Layer	Sequential Ops	Max Path Length
Self - Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$
Recurrent	$O(n \cdot d^2)$	$O(n)$	$O(n)$
Convolutional	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$
Self - Attention (restricted)	$O(r \cdot n \cdot d)$	$O(1)$	$O(n/r)$

Table 3: Time Complexity of some key components in YOLO (Table Form)

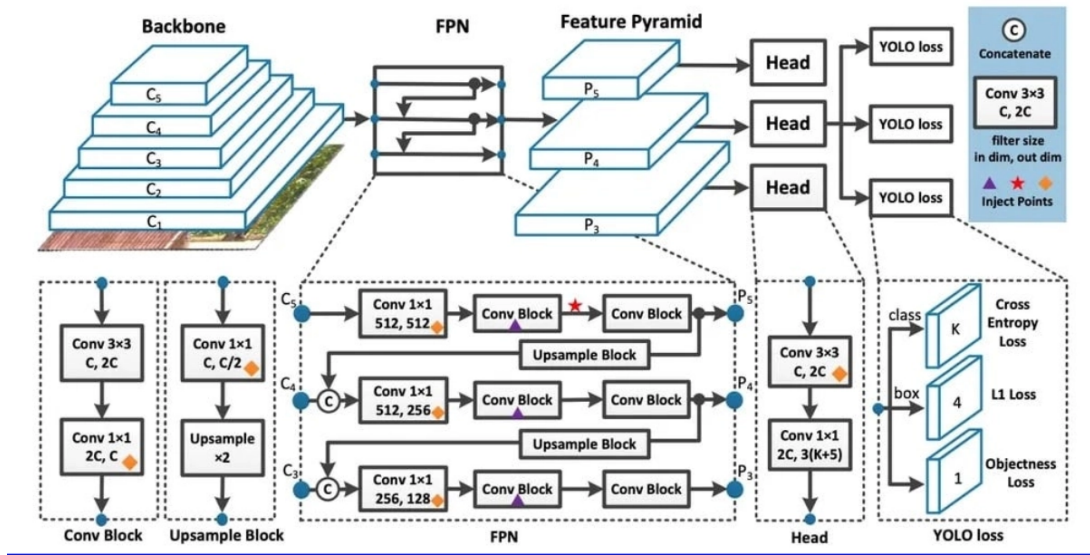


Figure 1: YOLO Architecture

ByteTrack Workflow

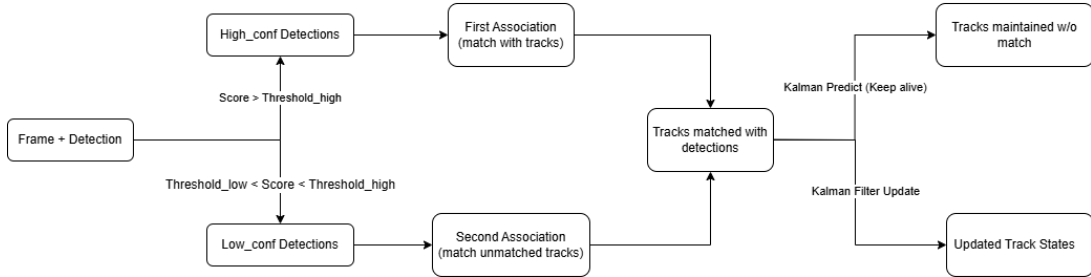


Figure 2: ByteTrack Workflow

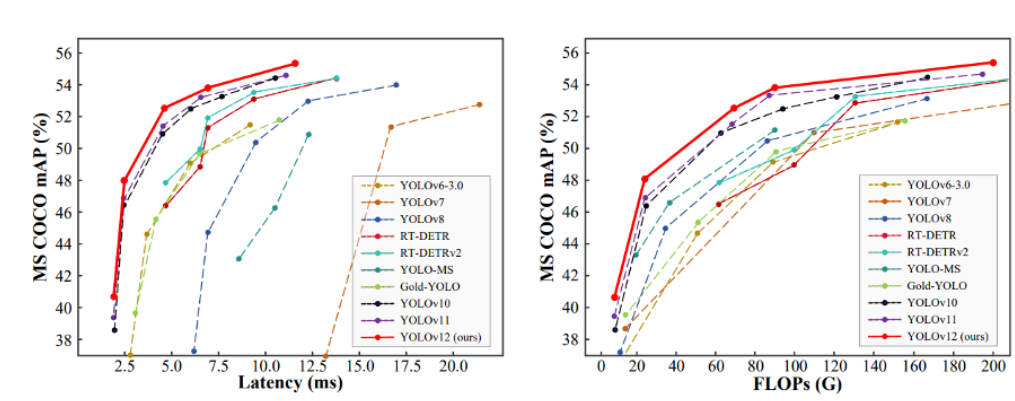


Figure 3: COCO Dataset

7 References

- Bochkovskiy, A., et al. (2020). YOLOv4. arXiv:2004.10934.
- Wang, C.Y., et al. (2022). YOLOv7. arXiv:2207.02696.
- Redmon, J., Divvala, S., Girshick, R., & Farhadi, A. (2016). You Only Look Once: Unified, Real-Time Object Detection. CVPR, 779–788.
- Szegedy, C., et al. (2015). Going Deeper with Convolutions. CVPR.
- Lin, M., Chen, Q., & Yan, S. (2013). Network in Network. arXiv:1312.4400.
- Terven, J., & Córdova-Esparza, D.-M. (2023). A Comprehensive Review of YOLO Architectures in Computer Vision: From YOLOv1 to YOLOv8 and YOLO-NAS. Machine Learning and Knowledge Extraction, 5(4), 1680–1716. <https://doi.org/10.3390/make5040083>
- Redmon, J., & Farhadi, A. (2017). YOLO9000: Better, Faster, Stronger. arXiv:1612.08242.
- Redmon, J., & Farhadi, A. (2018). YOLOv3: An Incremental Improvement. arXiv:1804.02767.
- Bochkovskiy, A., Wang, C.Y., & Liao, H.Y.M. (2020). YOLOv4: Optimal Speed and Accuracy of Object Detection. arXiv:2004.10934.
- Wang, C.Y., Bochkovskiy, A., & Liao, H.Y.M. (2022). YOLOv7: Trainable bag-of-freebies sets new state-of-the-art for real-time object detectors. arXiv:2207.02696.
- Jocher, G., et al. (2023). YOLOv8 by Ultralytics. <https://github.com/ultralytics/ultralytics>
- Kalman, R. E. (1960). "A New Approach to Linear Filtering and Prediction Problems." ASME. J. Basic Eng. March 1960; 82(1): 35–45. <https://doi.org/10.1115/1.3662552>
- Zhang, Y., Sun, P., Jiang, Y., Yu, D., Weng, F., Yuan, Z., Luo, P., Liu, W., & Wang, X. (2022). ByteTrack: Multi-object tracking by associating every detection box. arXiv. <https://arxiv.org/abs/2110.06864>
- Huang, C., Han, S., He, M., Zheng, W., & Wei, Y. (2024). DeconfuseTrack: Dealing with Confusion for Multi-Object Tracking. 2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 35, 19290–19299. <https://doi.org/10.1109/cvpr52733.2024.01825>
- Famili, A., Shen, W., Weber, R., & Simoudis, E. (1997). Data preprocessing and intelligent data analysis. Intelligent Data Analysis, 1(1–4), 3–23. [https://doi.org/10.1016/s1088-467x\(98\)00007-9](https://doi.org/10.1016/s1088-467x(98)00007-9)
- Kondrackis, L. (2024, September 5). How to Detect Small Objects: a guide. Roboflow Blog. <https://blog.roboflow.com/detect-small-objects/>
- Han, S., Wang, H., Yu, E., & Hu, Z. (2023). ORT: Occlusion-robust for multi-object tracking. Fundamental Research. <https://doi.org/10.1016/j.fmre.2023.02.003>
- Zhang, Y., Wang, T., & Zhang, X. (2023). MOTRV2: Bootstrapping End-to-End Multi-Object Tracking by pretrained Object Detectors. 2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR). <https://doi.org/10.1109/cvpr52729.2023.02112>